

## General Linear Model:

### 1. What is the purpose of the General Linear Model (GLM)?

The purpose of the General Linear Model (GLM) is to analyze and model the relationship between a dependent variable and one or more independent variables. It provides a framework for statistical techniques such as regression analysis, analysis of variance (ANOVA), and analysis of covariance (ANCOVA) to understand the impact of predictors on the outcome variable and make inferences about their relationships.

### 2. What are the key assumptions of the General Linear Model?

1. Linearity: The relationship between the dependent variable and independent variables is assumed to be linear.
2. Independence: The observations or data points used in the analysis are assumed to be independent of each other.
3. Normality: The residuals or errors of the model are assumed to follow a normal distribution.
4. Homoscedasticity: The variance of the residuals is assumed to be constant across all levels of the independent variables.
5. No multicollinearity: The independent variables used in the model should not be highly correlated with each other.
6. No influential outliers: The presence of influential outliers can disproportionately affect the model's results.

### 3. How do you interpret the coefficients in a GLM?

- GLM coefficients represent the estimated effects of independent variables on the dependent variable.
- Interpretation depends on the specific model type.
- In linear regression, a coefficient signifies the change in the dependent variable for a one-unit change in the independent variable, holding other variables constant.
- In logistic regression, coefficients represent changes in log-odds or odds ratio for a one-unit change in the predictor.
- For categorical variables, coefficients indicate differences in mean values between each category and a reference category.
- Consider the context, variable scales, and statistical significance for accurate interpretation.

### 4. What is the difference between a univariate and multivariate GLM?

- Univariate GLM analyzes a single dependent variable and examines its relationship with one or more independent variables.
- Multivariate GLM analyzes multiple dependent variables simultaneously and explores their relationships with one or more independent variables.
- Univariate GLM focuses on understanding the impact of predictors on a single outcome variable.
- Multivariate GLM allows for examining patterns, relationships, and differences across multiple outcome variables concurrently.
- The choice between univariate and multivariate GLM depends on the research question and the nature of the data being analyzed.

5. Explain the concept of interaction effects in a GLM.

Interaction effects in a GLM occur when the relationship between the dependent variable and one predictor depends on the level or condition of another predictor. They reveal that the combined impact of predictors is different from their individual effects. Interaction effects help understand how variables interact to influence the outcome. They are analyzed through interaction terms in the GLM equation and visualizations.

6. How do you handle categorical predictors in a GLM?

- Categorical predictors in a GLM are typically represented using dummy variables.
- Each category of the categorical predictor is encoded as a binary variable (0 or 1).
- The reference category is chosen as the baseline, and the other categories are compared to it.

7. What is the purpose of the design matrix in a GLM?

The design matrix in a GLM organizes the predictor variables into a matrix. It represents the relationship between the dependent variable and independent variables. Each column in the design matrix corresponds to a predictor variable, including continuous and categorical variables.

8. How do you test the significance of predictors in a GLM?

The significance of predictors in a GLM is typically tested using hypothesis tests such as t-tests or F-tests. The p-values associated with the tests indicate whether the predictor variables have a significant effect on the dependent variable. A smaller p-value suggests a stronger evidence of a significant relationship.

9. What is the difference between Type I, Type II, and Type III sums of squares in a GLM?

- Type I, Type II, and Type III sums of squares refer to different methods of partitioning the sum of squares in a GLM.
- Type I sums of squares focus on the unique contribution of each predictor, sequentially adding variables to the model.
- Type II sums of squares consider the unique contribution of each predictor while controlling for other predictors in the model.
- Type III sums of squares examine the contribution of each predictor, independent of the order in which variables are entered into the model.

10. Explain the concept of deviance in a GLM.

- Deviance is a measure of the lack of fit between the observed data and the fitted model in a GLM.
- It quantifies the discrepancy between the predicted values and the actual values of the dependent variable.
- Lower deviance indicates a better fit of the model to the data.
- Deviance is used to compare different models and assess their goodness of fit.

## **Regression:**

11. What is regression analysis and what is its purpose?

- Regression analysis is a statistical technique used to model the relationship between a dependent variable and one or more independent variables.
- Its purpose is to understand and quantify the impact of the independent variables on the dependent variable, make predictions, and identify patterns and trends in the data.

12. What is the difference between simple linear regression and multiple linear regression?

- Simple linear regression involves one dependent variable and one independent variable, capturing a linear relationship between them.
- Multiple linear regression involves one dependent variable and two or more independent variables, allowing for the analysis of multiple predictors simultaneously.

13. How do you interpret the R-squared value in regression?

- The R-squared value represents the proportion of the variance in the dependent variable that can be explained by the independent variables.
- It indicates the goodness of fit of the regression model: a higher R-squared value suggests a better fit, with more variance explained by the predictors.

14. What is the difference between correlation and regression?

- Correlation measures the strength and direction of the linear relationship between two variables, focusing on their association.
- Regression analysis goes beyond correlation by modeling the relationship between the dependent variable and one or more independent variables, enabling prediction and understanding of the impact of predictors on the outcome.

15. What is the difference between the coefficients and the intercept in regression?

- Coefficients (slopes) in regression represent the change in the dependent variable for a one-unit change in the corresponding independent variable, while holding other variables constant.
- The intercept is the value of the dependent variable when all independent variables are zero, providing the starting point for the regression line.

16. How do you handle outliers in regression analysis?

- Outliers can have a significant impact on regression analysis. Handling them may involve:
  - Investigating the outliers to determine their validity and potential causes.
  - Transforming the data or using robust regression techniques that are less sensitive to outliers.
  - Considering removing or adjusting outliers, if they are determined to be data errors or have undue influence on the model.

17. What is the difference between ridge regression and ordinary least squares regression?

- Ordinary Least Squares (OLS) regression aims to minimize the sum of squared residuals and estimate coefficients without constraints.
- Ridge regression adds a penalty term to the OLS objective function to shrink the coefficient estimates, reducing their sensitivity to multicollinearity and providing better stability.

18. What is heteroscedasticity in regression and how does it affect the model?

- Heteroscedasticity refers to the unequal spread of residuals across the range of predictor variables.
- It violates the assumption of homoscedasticity in regression, potentially affecting the validity of statistical tests and confidence intervals.
- To address heteroscedasticity, robust standard errors or transformations (e.g., log transformations) can be used, or weighted regression models can be employed.

19. How do you handle multicollinearity in regression analysis?

- Multicollinearity occurs when independent variables are highly correlated with each other.
- To handle multicollinearity, approaches include:
  - Dropping one of the highly correlated variables.
  - Combining or transforming the correlated variables.
  - Using dimensionality reduction techniques, such as principal component analysis (PCA) or factor analysis.

20. What is polynomial regression and when is it used?

- Polynomial regression is a form of regression analysis where the relationship between the independent variable(s) and dependent variable is modeled as an  $n$ th-degree polynomial.
- It is used when there is a nonlinear relationship between the variables, allowing for curved or nonlinear patterns to be captured in the regression model.

## **Loss function:**

21. What is a loss function and what is its purpose in machine learning?

- A loss function measures the discrepancy between predicted and actual values in machine learning.
- Its purpose is to quantify the model's performance and guide the learning process by providing a measure of how well the model is fitting the data.

22. What is the difference between a convex and non-convex loss function?

- A convex loss function has a single global minimum, making optimization easier.
- A non-convex loss function has multiple local minima, making optimization more challenging.

23. What is mean squared error (MSE) and how is it calculated?

- Mean squared error (MSE) measures the average squared difference between predicted and actual values. It is calculated by summing the squared differences and dividing by the number of data points.

24. What is mean absolute error (MAE) and how is it calculated?

- Mean absolute error (MAE) measures the average absolute difference between predicted and actual values.
- It is calculated by summing the absolute differences and dividing by the number of data points.

25. What is log loss (cross-entropy loss) and how is it calculated?

- Log loss, also known as cross-entropy loss, is used for binary or multi-class classification problems.
- It quantifies the difference between predicted probabilities and actual class labels.
- It is calculated by taking the negative logarithm of the predicted probabilities of the true class.

26. How do you choose the appropriate loss function for a given problem?

- The choice of loss function depends on the nature of the problem and the desired behavior of the model.
- For regression problems, MSE or MAE are commonly used.
- For classification problems, log loss or other classification-specific loss functions (e.g., hinge loss, softmax loss) may be appropriate.

27. Explain the concept of regularization in the context of loss functions.

- Regularization is a technique used to prevent overfitting by adding a penalty term to the loss function.
- It discourages complex models and encourages simpler models by penalizing large coefficients or model complexity.
- Common regularization techniques include L1 regularization (Lasso) and L2 regularization (Ridge).

28. What is Huber loss and how does it handle outliers?

- Huber loss is a loss function that is less sensitive to outliers compared to squared loss (MSE).
- It combines squared loss for small errors and absolute loss for large errors.
- Huber loss is a compromise between squared loss and absolute loss, providing robustness to outliers.

29. What is quantile loss and when is it used?

- Quantile loss measures the discrepancies between predicted and actual quantiles of a distribution.
- It is used when the focus is on estimating a specific quantile or when asymmetric errors are more important than overall accuracy.

30. What is the difference between squared loss and absolute loss?

- Squared loss penalizes larger errors more than absolute loss due to squaring the differences.
- Absolute loss treats all errors equally and is less sensitive to outliers.
- Squared loss has better mathematical properties for optimization, while absolute loss is more robust to extreme values.

## Optimizer (GD):

31. What is an optimizer and what is its purpose in machine learning?

- An optimizer is an algorithm used to minimize the loss function and optimize the model parameters during the training process.
- Its purpose is to iteratively update the model parameters to find the optimal values that minimize the discrepancy between predicted and actual values.

32. What is Gradient Descent (GD) and how does it work?

- Gradient Descent is an optimization algorithm used to find the minimum of a function.
- It starts with an initial guess for the parameters and iteratively updates them in the opposite direction of the gradient of the loss function.
- By following the negative gradient, it aims to reach the minimum of the function.

33. What are the different variations of Gradient Descent?

- Different variations of Gradient Descent include Batch Gradient Descent, Stochastic Gradient Descent, and Mini-Batch Gradient Descent.
- They differ in the amount of data used to compute the gradient and the update strategy.

34. What is the learning rate in GD and how do you choose an appropriate value?

- The learning rate controls the step size taken during each parameter update in GD.
- Choosing an appropriate learning rate is crucial; a high value may cause divergence, while a low value may result in slow convergence.
- It is often determined through experimentation and validation, starting with a small value and gradually increasing it until satisfactory results are achieved.

35. How does GD handle local optima in optimization problems?

- GD can get trapped in local optima, where the loss function is relatively small compared to the immediate surrounding area.
- To mitigate this, various techniques can be employed, such as using different initial parameter values, trying different optimization algorithms, or incorporating regularization.

36. What is Stochastic Gradient Descent (SGD) and how does it differ from GD?

- Stochastic Gradient Descent updates the parameters using the gradient computed on a single training example at each iteration.
- Unlike GD, which uses the entire training set, SGD is computationally more efficient but introduces more noise and variance in the parameter updates.

37. Explain the concept of batch size in GD and its impact on training.

- Batch size refers to the number of training examples used to compute the gradient and update the parameters in each iteration of GD.
- A larger batch size provides a more accurate estimate of the gradient but increases memory requirements and slows down computation.
- A smaller batch size introduces more stochasticity but can make the optimization process faster.

38. What is the role of momentum in optimization algorithms?

- Momentum is a term that adds inertia to the optimization process.
- It allows the optimizer to keep track of previous parameter updates and smooths out fluctuations in the gradient direction.
- This helps accelerate convergence, especially in the presence of noise or high curvature in the loss landscape.

39. What is the difference between batch GD, mini-batch GD, and SGD?

- Batch Gradient Descent uses the entire training set to compute the gradient and update the parameters.
- Mini-Batch Gradient Descent uses a subset (batch) of training examples to compute the gradient and update the parameters.
- Stochastic Gradient Descent uses a single training example to compute the gradient and update the parameters.

40. How does the learning rate affect the convergence of GD?

- The learning rate determines the step size taken in each parameter update.
- A large learning rate can lead to overshooting the optimal values and divergence.
- A small learning rate may cause slow convergence.
- The learning rate needs to be carefully chosen to balance convergence speed and stability.

## **Regularization:**

41. What is regularization and why is it used in machine learning?

- Regularization is a technique used to prevent overfitting and improve the generalization of machine learning models.
- It adds a penalty term to the loss function, discouraging complex models and reducing the impact of certain features.

42. What is the difference between L1 and L2 regularization?

- L1 regularization (Lasso) adds the sum of absolute values of the coefficients as a penalty term.
- L2 regularization (Ridge) adds the sum of squared values of the coefficients as a penalty term.
- L1 regularization promotes sparsity and feature selection, while L2 regularization encourages small, non-zero coefficients.

43. Explain the concept of ridge regression and its role in regularization.

- Ridge regression is a type of linear regression that incorporates L2 regularization.
- It adds the sum of squared coefficients to the loss function, shrinking their values and reducing their impact on the model.
- Ridge regression helps mitigate the effects of multicollinearity and provides more stable coefficient estimates.

44. What is the elastic net regularization and how does it combine L1 and L2 penalties?

- Elastic Net regularization combines both L1 and L2 penalties in the loss function.
- It adds a linear combination of the L1 (Lasso) and L2 (Ridge) penalties to strike a balance between sparsity and shrinkage.
- Elastic Net is useful when dealing with highly correlated features and when feature selection is desired.

45. How does regularization help prevent overfitting in machine learning models?

- Regularization helps prevent overfitting by reducing the complexity and flexibility of a model.
- It discourages large coefficient values, suppressing noise and reducing the model's sensitivity to specific data points.
- Regularization encourages the model to learn more general patterns and prevents it from memorizing the training data too closely.

46. What is early stopping and how does it relate to regularization?

- Early stopping is a technique used to prevent overfitting by monitoring the model's performance on a validation set during training.
- It stops the training process when the performance on the validation set starts to degrade, preventing the model from learning the noise in the training data.
- Early stopping can be seen as a form of regularization, as it limits the model's capacity to fit the training data too closely.

47. Explain the concept of dropout regularization in neural networks.

- Dropout regularization is a technique used in neural networks to prevent overfitting.
- It randomly "drops out" a fraction of the neurons during each training iteration, forcing the network to learn more robust and generalizable representations.
- Dropout acts as a form of ensemble learning, as the network learns to adapt with different sets of neurons active at each iteration.

48. How do you choose the regularization parameter in a model?

- The choice of the regularization parameter depends on the specific model and dataset.
- It is typically determined through cross-validation, trying different values and evaluating the model's performance on validation data.
- The optimal regularization parameter balances the trade-off between model complexity and generalization performance.

49. What is the difference between feature selection and regularization?

- Feature selection aims to select a subset of relevant features to improve model performance and interpretability.
- Regularization, on the other hand, applies a penalty to the model's coefficients, encouraging simpler models and reducing the impact of irrelevant features.
- Feature selection explicitly removes features, while regularization shrinks the coefficients towards zero, making them less influential.



50. What is the trade-off between bias and variance in regularized models?

- Regularized models strike a trade-off between bias and variance.
- By adding a regularization penalty, the model reduces its variance by reducing the impact of individual features and decreasing overfitting.
- However, this regularization can introduce a small amount of bias, as the model may not fit the training data as closely as an unregularized model.

## **SVM:**

51. What is Support Vector Machines (SVM) and how does it work?

- Support Vector Machines (SVM) is a supervised machine learning algorithm used for classification and regression tasks.
- SVM finds the optimal hyperplane that maximally separates the data points of different classes by considering the support vectors, which are the data points closest to the decision boundary.

52. How does the kernel trick work in SVM?

- The kernel trick in SVM allows the algorithm to implicitly transform the input data into a higher-dimensional feature space without explicitly computing the transformations.
- It enables SVM to handle non-linearly separable data by projecting it into a higher-dimensional space where linear separation becomes possible.

53. What are support vectors in SVM and why are they important?

- Support vectors are the data points closest to the decision boundary of an SVM model.
- They play a crucial role in defining the decision boundary and determining the model's performance.
- Support vectors have non-zero coefficients and are essential for making predictions on new data points.

54. Explain the concept of the margin in SVM and its impact on model performance.

- The margin in SVM refers to the distance between the decision boundary and the nearest data points from each class.
- A larger margin indicates a more robust and better-generalized model, as it provides a wider separation between classes and reduces the risk of misclassification.

55. How do you handle unbalanced datasets in SVM?

- Handling unbalanced datasets in SVM can involve techniques such as:
  - Adjusting the class weights to give more importance to the minority class.
  - Oversampling the minority class or undersampling the majority class to balance the class distribution.
  - Using techniques like SMOTE (Synthetic Minority Over-sampling Technique) to generate synthetic examples of the minority class.

56. What is the difference between linear SVM and non-linear SVM?

- Linear SVM uses a linear decision boundary to separate the classes.
- Non-linear SVM uses a kernel function to map the data into a higher-dimensional space, where a linear decision boundary can be found, allowing for non-linear separation.

57. What is the role of C-parameter in SVM and how does it affect the decision boundary?

- The C-parameter in SVM controls the trade-off between maximizing the margin and minimizing the training errors.
- A smaller C-value leads to a wider margin but allows more training errors, potentially resulting in a more general model.
- A larger C-value focuses on minimizing training errors, potentially resulting in a narrower margin and higher complexity.

58. Explain the concept of slack variables in SVM.

- Slack variables are introduced in soft margin SVM to allow for some misclassification of data points.
- They measure the degree to which data points violate the margin or fall on the wrong side of the decision boundary.
- Slack variables provide flexibility in handling non-linearly separable data or cases with outliers.

59. What is the difference between hard margin and soft margin in SVM?

- Hard margin SVM aims to find a decision boundary that perfectly separates the data without allowing any misclassifications.
- Soft margin SVM introduces a margin of tolerance and allows for some misclassifications to handle non-linearly separable data or noisy datasets.

60. How do you interpret the coefficients in an SVM model?

- In linear SVM, the coefficients represent the weights assigned to the features.
- Positive coefficients indicate a positive influence on the classification decision, while negative coefficients indicate a negative influence.
- The magnitude of the coefficients indicates the importance of the corresponding feature in the decision-making process.

## **Decision Trees:**

61. What is a decision tree and how does it work?

- A decision tree is a supervised machine learning algorithm used for classification and regression tasks.
- It works by recursively partitioning the feature space based on a set of decision rules to create a tree-like structure that predicts the target variable.

62. How do you make splits in a decision tree?

- Splits in a decision tree are made based on the feature values that best separate the data according to a certain criterion.

- The algorithm searches for the best split by evaluating different features and splitting points based on specific impurity measures or information gain.

63. What are impurity measures (e.g., Gini index, entropy) and how are they used in decision trees?

- Impurity measures quantify the disorder or impurity of a set of samples within a node in a decision tree.
- Gini index and entropy are two commonly used impurity measures in decision trees.
- They are used to assess the quality of a split and determine the optimal feature and splitting point to maximize the separation of classes or the reduction of uncertainty.

64. Explain the concept of information gain in decision trees.

- Information gain measures the reduction in entropy or impurity achieved by splitting the data based on a particular feature.
- It quantifies the amount of information gained about the target variable after making a split, helping to determine the most informative feature for the split.

65. How do you handle missing values in decision trees?

- Missing values in decision trees can be handled by either ignoring the samples with missing values, treating them as a separate category, or imputing the missing values based on other observations.
- The decision tree algorithm can be modified to handle missing values during the split evaluation process.

66. What is pruning in decision trees and why is it important?

- Pruning is a technique used to reduce the complexity of decision trees by removing unnecessary branches or nodes.
- It helps prevent overfitting and improves the model's ability to generalize to unseen data.
- Pruning can be performed based on criteria such as the reduction in overall impurity or the evaluation of the model's performance on validation data.

67. What is the difference between a classification tree and a regression tree?

- A classification tree is used for categorical or discrete target variables, where the goal is to classify the data into different classes or categories.
- A regression tree is used for continuous target variables, where the goal is to predict a numerical value or estimate a function.

68. How do you interpret the decision boundaries in a decision tree?

- Decision boundaries in a decision tree are represented by the splits and branches of the tree structure.
- Each split creates a partition in the feature space that separates the data into different regions or classes based on the decision rules.
- The decision boundaries can be interpreted as the rules or conditions that determine the class assignment or prediction.

69. What is the role of feature importance in decision trees?

- Feature importance measures the relative contribution or importance of each feature in the decision-making process of a decision tree.
- It helps identify the most influential features and understand their impact on the model's predictions.
- Feature importance can be derived from metrics such as the Gini importance or the average depth of features in the tree.

70. What are ensemble techniques and how are they related to decision trees?

- Ensemble techniques combine multiple individual models, such as decision trees, to improve predictive performance.
- They leverage the diversity of the individual models to enhance generalization and reduce overfitting.
- Popular ensemble techniques that use decision trees as base models include Random Forest, Gradient Boosting, and AdaBoost.

## **Ensemble Techniques:**

71. What are ensemble techniques in machine learning?

- Ensemble techniques in machine learning combine multiple individual models to make predictions or classifications.
- They aim to leverage the collective knowledge and diversity of the models to improve overall performance and enhance generalization.

72. What is bagging and how is it used in ensemble learning?

- Bagging (Bootstrap Aggregating) is an ensemble technique that involves training multiple models on different subsets of the training data and combining their predictions.
- It helps reduce variance and overfitting by averaging or voting the predictions of individual models.

73. Explain the concept of bootstrapping in bagging.

- Bootstrapping is a resampling technique used in bagging where multiple subsets of the training data are created by random sampling with replacement.
- Each subset, known as a bootstrap sample, has the same size as the original training data but may contain duplicate instances.
- These bootstrap samples are then used to train individual models in the ensemble.

74. What is boosting and how does it work?

- Boosting is an ensemble technique that iteratively trains weak models and gives more weight to misclassified instances.
- It focuses on samples that are harder to classify, aiming to improve the overall model's performance.
- The predictions of the weak models are combined to create a strong, boosted model.

75. What is the difference between AdaBoost and Gradient Boosting?

- AdaBoost (Adaptive Boosting) and Gradient Boosting are both boosting algorithms but differ in how they assign weights to instances and how they update the model in each iteration.
- AdaBoost adjusts the instance weights to focus on misclassified instances, while Gradient Boosting fits subsequent models to the residuals of the previous models.

76. What is the purpose of random forests in ensemble learning?

- Random forests are an ensemble technique that combines multiple decision trees in a parallel manner.
- They reduce overfitting, improve generalization, and handle high-dimensional data by using random subsets of features for each tree and averaging their predictions.

77. How do random forests handle feature importance?

- Random forests determine feature importance based on how much each feature reduces the impurity or error in the model.
- The importance of a feature is calculated by averaging the impurity decrease or the reduction in the model's error caused by that feature across all trees in the forest.

78. What is stacking in ensemble learning and how does it work?

- Stacking is an ensemble technique that combines the predictions of multiple models as input to a meta-model, which learns to make the final prediction.
- The predictions of the individual models serve as additional features for the meta-model, enabling it to learn a more sophisticated combination of predictions.

79. What are the advantages and disadvantages of ensemble techniques?

- Advantages:
  - Improved predictive performance and generalization.
  - Ability to handle complex patterns and non-linear relationships.
  - Increased stability and robustness to noisy data.
- Disadvantages:
  - Increased computational complexity and training time.
  - Lack of interpretability compared to individual models.
  - Potential overfitting if the ensemble is too complex or diverse.

80. How do you choose the optimal number of models in an ensemble?

- The optimal number of models in an ensemble depends on the specific problem, dataset, and performance evaluation.
- It is often determined through cross-validation or validation set performance, where the performance stabilizes or starts to degrade after reaching an optimal number of models.
- Increasing the number of models beyond this point may not provide significant improvement and can lead to increased complexity and computational cost.